

ECE 432/532  
Programming for Parallel Processors

# Sorting

- $n$  keys and  $p = \text{comm sz processes}$ .
- $n/p$  keys assigned to each process.
- No restrictions on which keys are assigned to which processes.
- When the algorithm terminates:
  - The keys assigned to each process should be sorted in (say) increasing order.
  - If  $0 \leq q < r < p$ , then each key assigned to process  $q$  should be less than or equal to every key assigned to process  $r$ .

# Serial bubble sort

```
void Bubble_sort(  
    int a[] /* in/out */,  
    int n /* in */) {  
    int list_length, i, temp;  
  
    for (list_length = n; list_length >= 2; list_length--)  
        for (i = 0; i < list_length-1; i++)  
            if (a[i] > a[i+1]) {  
                temp = a[i];  
                a[i] = a[i+1];  
                a[i+1] = temp;  
            }  
  
} /* Bubble_sort */
```

6 5 3 1 8 7 2 4

# Serial bubble sort

- Is it good for parallel implementation?
- Example, you have the numbers 9, 5, 7
  - How would you sort them in serial way?
  - How would you sort them in parallel way?
- The order in which the “compare-swaps” take place is essential to the correctness of the algorithm

# Odd-even transposition sort

- A sequence of phases.
- Even phases, compare swaps:

$$(a[0], a[1]), (a[2], a[3]), (a[4], a[5]), \dots$$

- Odd phases, compare swaps:

$$(a[1], a[2]), (a[3], a[4]), (a[5], a[6]), \dots$$

# Example

Start: 5, 9, 4, 3


Even phase: compare-swap (5,9) and (4,3)  
getting the list 5, 9, 3, 4

Odd phase: compare-swap (9,3)  
getting the list 5, 3, 9, 4

Even phase: compare-swap (5,3) and (9,4)  
getting the list 3, 5, 4, 9

Odd phase: compare-swap (5,4)  
getting the list 3, 4, 5, 9

# Example



# Serial odd-even transposition sort

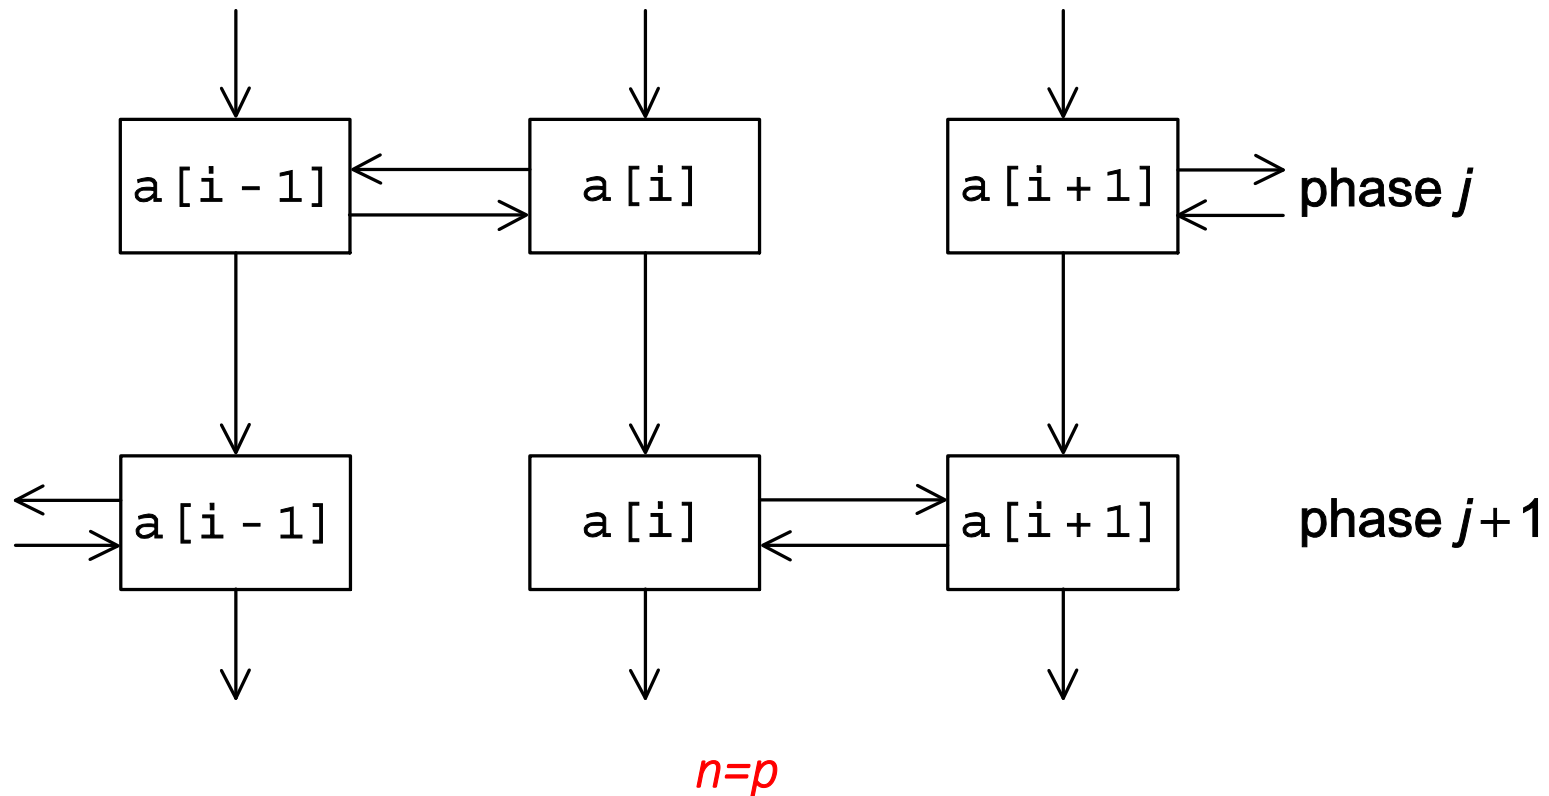
```
void Odd_even_sort(  
    int a[] /* in/out */,  
    int n /* in */) {  
    int phase, i, temp;  
  
    for (phase = 0; phase < n; phase++)  
        if (phase % 2 == 0) { /* Even phase */  
            for (i = 1; i < n; i += 2)  
                if (a[i-1] > a[i]) {  
                    temp = a[i];  
                    a[i] = a[i-1];  
                    a[i-1] = temp;  
                }  
        } else { /* Odd phase */  
            for (i = 1; i < n-1; i += 2)  
                if (a[i] > a[i+1]) {  
                    temp = a[i];  
                    a[i] = a[i+1];  
                    a[i+1] = temp;  
                }  
        }  
    } /* Odd_even_sort */
```



# Parallel odd-even transposition sort

- The odd-even transposition sort has considerably more opportunities for parallelism than bubble sort
- All of the compare-swaps in a single phase can happen simultaneously.
- Foster's methodology:
  - *Tasks*: Determine the value of  $a[i]$  at the end of phase  $j$ .
  - *Communications*: The task that's determining the value of  $a[i]$  needs to communicate with either the task determining the value of  $a[i-1]$  or  $a[i+1]$ . Also the value of  $a[i]$  at the end of phase  $j$  needs to be available for determining the value of  $a[i]$  at the end of phase  $j+1$ .

# Communications among tasks in odd-even sort



# Parallel odd-even transposition sort

| Time             | Process       |                |                |                |
|------------------|---------------|----------------|----------------|----------------|
|                  | 0             | 1              | 2              | 3              |
| Start            | 15, 11, 9, 16 | 3, 14, 8, 7    | 4, 6, 12, 10   | 5, 2, 13, 1    |
| After Local Sort | 9, 11, 15, 16 | 3, 7, 8, 14    | 4, 6, 10, 12   | 1, 2, 5, 13    |
| After Phase 0    | 3, 7, 8, 9    | 11, 14, 15, 16 | 1, 2, 4, 5     | 6, 10, 12, 13  |
| After Phase 1    | 3, 7, 8, 9    | 1, 2, 4, 5     | 11, 14, 15, 16 | 6, 10, 12, 13  |
| After Phase 2    | 1, 2, 3, 4    | 5, 7, 8, 9     | 6, 10, 11, 12  | 13, 14, 15, 16 |
| After Phase 3    | 1, 2, 3, 4    | 5, 6, 7, 8     | 9, 10, 11, 12  | 13, 14, 15, 16 |

When will the list be sorted? After how many phases?

# Pseudo-code

```
Sort local keys;
for (phase = 0; phase < comm_sz; phase++) {
    partner = Compute_partner(phase, my_rank);
    if (I'm not idle) {
        Send my keys to partner;
        Receive keys from partner;
        if (my_rank < partner)
            Keep smaller keys;
        else
            Keep larger keys;
    }
}
```

# Compute\_partner

```
if (phase % 2 == 0)          /* Even phase */
    if (my_rank % 2 != 0)     /* Odd rank */
        partner = my_rank - 1;
    else                      /* Even rank */
        partner = my_rank + 1;
else                          /* Odd phase */
    if (my_rank % 2 != 0)     /* Odd rank */
        partner = my_rank + 1;
    else                      /* Even rank */
        partner = my_rank - 1;
if (partner == -1 || partner == comm_sz)
    partner = MPI_PROC_NULL;
```

MPI PROC NULL is a constant defined by MPI. When it's used as the source or destination rank in a point-to-point communication, no communication will take place and the call to the communication will simply return