




ECE 432/532
Programming for Parallel Processors

More on OpenMP directives

- **Work-Sharing Constructs**


- A work-sharing construct divides the execution of the enclosed code region among the members of the team that encounter it.
- Work-sharing constructs do not launch new threads
- There is no implied barrier upon entry to a work-sharing construct, however there is an implied barrier at the end of a work sharing construct.

More on OpenMP directives



Work-Sharing Constructs

- **DO / for** - shares iterations of a loop across the team
- Represents a type of "data parallelism"



```
#include <omp.h>

#define N 1000
#define CHUNKSIZE 100

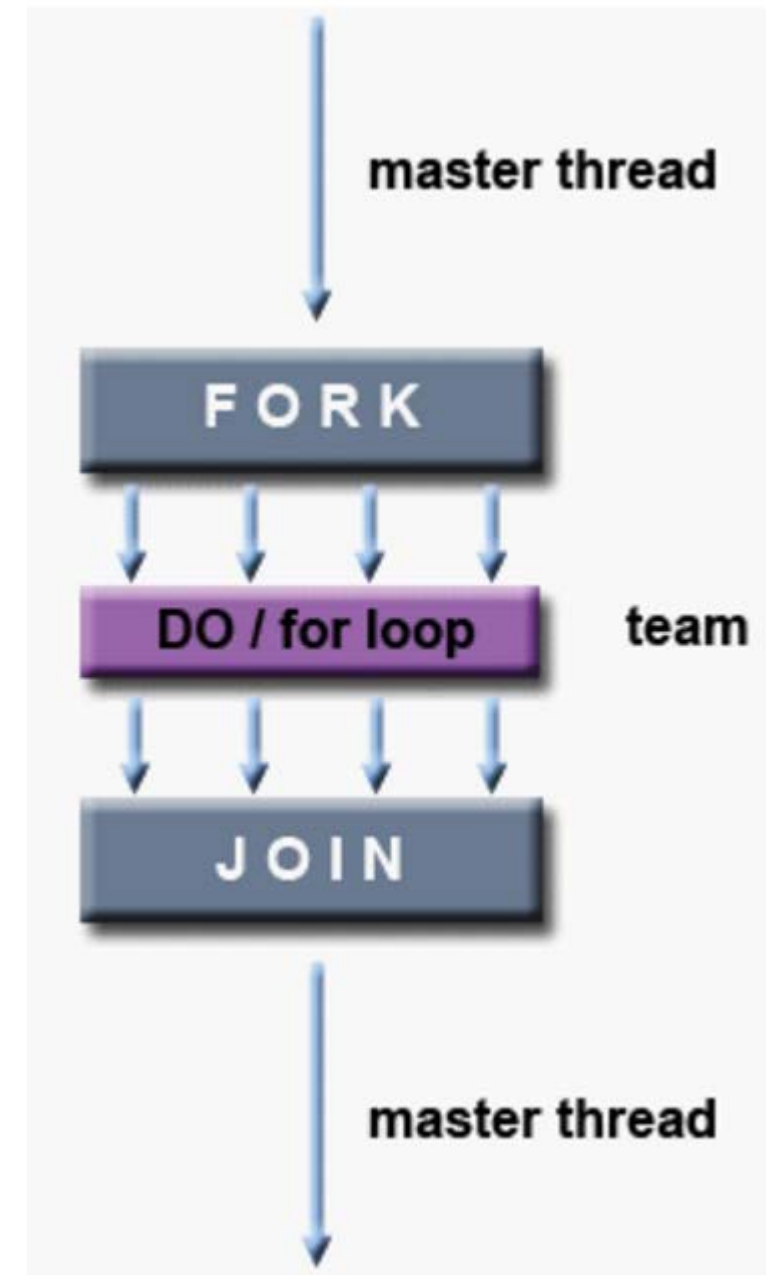
main(int argc, char *argv[]) {
    int i, chunk;
    float a[N], b[N], c[N];

    for (i=0; i < N; i++)
        a[i] = b[i] = i * 1.0;
    chunk = CHUNKSIZE;

    #pragma omp parallel shared(a,b,c,chunk) private(i)
    {
        #pragma omp for schedule(dynamic,chunk) nowait
        for (i=0; i < N; i++)
            c[i] = a[i] + b[i];
    } /* end of parallel region */
}
```

Work-Sharing Constructs

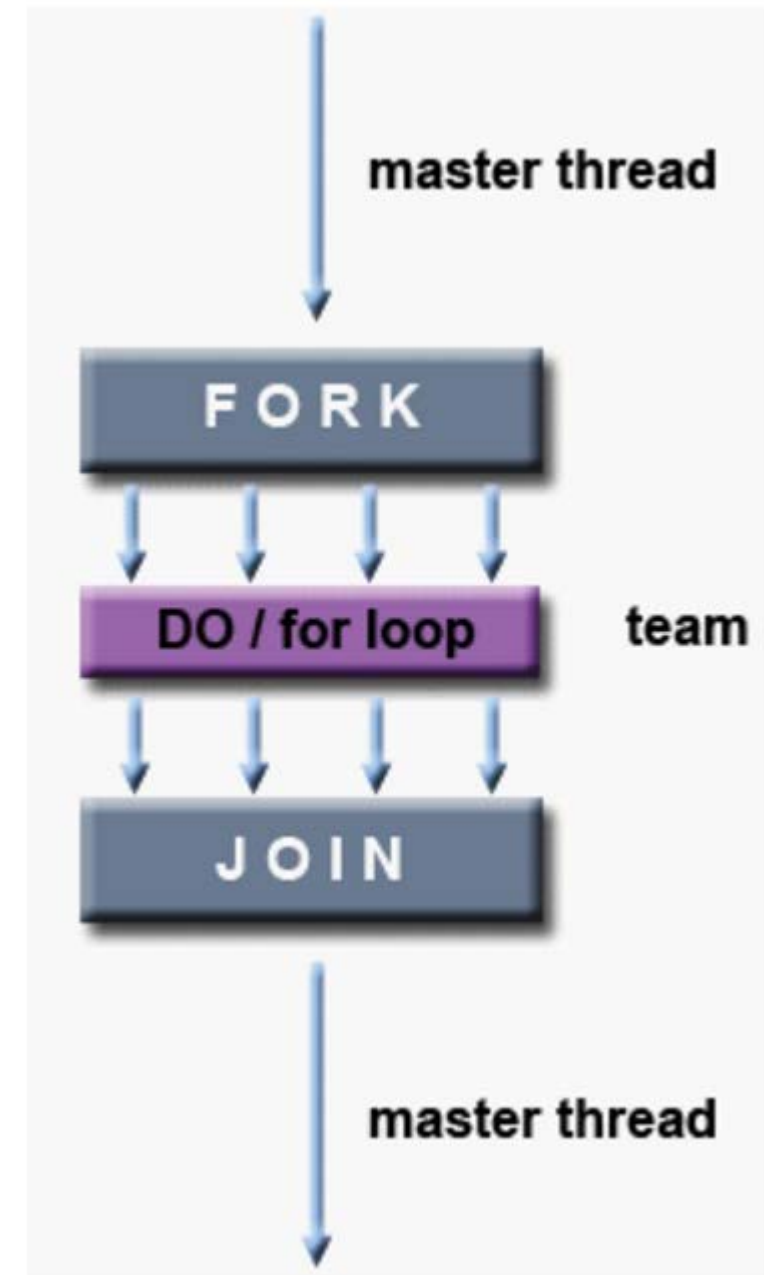
- The *for* directive specifies that the iterations of the loop immediately following it must be executed in parallel by the team
- This assumes a parallel region has already been initiated, otherwise it executes in serial on a single processor.



Work-Sharing Constructs


- **Restrictions:**

- The loop has to be a for loop
- The loop iteration variable must be an integer and the loop control parameters must be the same for all threads
- Program correctness must not depend upon which thread executes a particular iteration
- It is illegal to branch out of a loop associated with a *for* directive
- The chunk size must be specified as a loop invariant integer expression, as there is no synchronization during its evaluation by different threads




Work-Sharing Constructs

- **SECTIONS** - breaks work into separate, discrete sections.
- Each section is executed by a thread.
- Can be used to implement a type of "functional parallelism".



Work-Sharing Constructs

- The **SECTIONS** directive is a non-iterative work-sharing construct.
 - It specifies that the enclosed section(s) of code are to be divided among the threads in the team.
- Independent **SECTION** directives are nested within a **SECTIONS** directive
 - Each **SECTION** is executed once by a thread in the team.
 - Different sections may be executed by different threads.
 - It is possible for a thread to execute more than one section



```
main(int argc, char *argv[]) {
    int i;
    float a[N], b[N], c[N], d[N];


    for (i=0; i < N; i++) {
        a[i] = i * 1.5;
        b[i] = i + 22.35;
    }

    #pragma omp parallel shared(a,b,c,d) private(i)
    {
        #pragma omp sections nowait
        {
            #pragma omp section
            for (i=0; i < N; i++)
                c[i] = a[i] + b[i];

            #pragma omp section
            for (i=0; i < N; i++)
                d[i] = a[i] * b[i];
        } /* end of sections */
    } /* end of parallel region */
}
```


Work-Sharing Constructs

- There is an implied barrier at the end of a **SECTIONS** directive, unless the **nowait** clause is used
- Q1: What happens if the number of threads and the number of SECTIONS are different? More threads than SECTIONS? Less threads than SECTIONS?
- Q2: Which thread executes which SECTION?



Work-Sharing Constructs

- **SINGLE** - serializes a section of code
- The **SINGLE** directive specifies that the enclosed code is to be executed by only one thread in the team
- May be useful when dealing with sections of code that are not thread safe (such as I/O)
- Threads in the team that do not execute the **SINGLE** directive, wait at the end of the enclosed code block, unless a **NOWAIT/nowait** clause is specified



```
#include <stdio.h>
#include <omp.h>

int main() {
    #pragma omp parallel num_threads(2)
    {
        #pragma omp single
        // Only a single thread can read the input.
        printf("read input\n");

        // Multiple threads in the team compute the results.
        printf("compute results\n");

        #pragma omp single
        // Only a single thread can write the output.
        printf("write output\n");
    }
}
```

Example - Is it correct?

```
#pragma omp parallel shared(n,a,b,c,d,sum) private(i)
{
    #pragma omp for nowait
    for (i=0; i<n; i++)
        a[i] += b[i];

    #pragma omp for nowait
    for (i=0; i<n; i++)
        c[i] += d[i];
    #pragma omp barrier

    #pragma omp for nowait reduction(+:sum)
    for (i=0; i<n; i++)
        sum += a[i] + c[i];
} /*-- End of parallel region --*/
```

Other directives

- ORDERED
- FIRSTPRIVATE
- LASTPRIVATE
- MASTER

ORDERED

- Specifies that code under a parallelized for loop should be executed like a sequential loop
 - The ordered directive must be within the dynamic extent of a for or parallel for construct with an ordered clause
 - The ordered directive supports no OpenMP clauses.

#pragma omp ordered
structured-block


```
#include <stdio.h>
#include <omp.h>

static float a[1000], b[1000], c[1000];

void test(int first, int last){
    #pragma omp for schedule(static) ordered
    for (int i = first; i <= last; ++i) {
        // Do something here.
        if (i % 2) {
            #pragma omp ordered
            printf("test() iteration %d\n", i);
        }
    }
}

void test2(int iter) {
    #pragma omp ordered
    printf("test2() iteration %d\n", iter);
}
```

```
int main( ) {  
    int i;  
    #pragma omp parallel  
    {  
        test(1, 8);  
        #pragma omp for ordered  
        for (i = 0 ; i < 5 ; i++)  
            test2(i);  
    }  
}
```

```
int main( ) {  
    int i;  
    #pragma omp parallel  
    {  
        test(1, 8);  
        #pragma omp for ordered  
        for (i = 0 ; i < 5 ; i++)  
            test2(i);  
    }  
}
```

test() iteration 1
test() iteration 3
test() iteration 5
test() iteration 7
test2() iteration 0
test2() iteration 1
test2() iteration 2
test2() iteration 3
test2() iteration 4

FIRSTPRIVATE

- Specifies that each thread should have its own instance of a variable, and that the variable should be initialized with the value of the variable, because it exists before the parallel construct.

firstprivate(var)

LASTPRIVATE

- Specifies that the enclosing context's version of the variable is set equal to the private version of whichever thread executes the final iteration (for-loop construct) or last section (#pragma sections).

lastprivate(var)

```
j = jstart;  
#pragma omp parallel for firstprivate(j)  
{  
    for(i=1; i<=n; i++){  
        if(i == 1 || i == n)  
            j = j + 1;  
        a[i] = a[i] + j;  
    }  
}
```

```
#pragma omp parallel for lastprivate(x)  
{  
    for(i=1; i<=n; i++){  
        x = sin( pi * dx * (float)i );  
        a[i] = exp(x);  
    }  
}  
lastx = x;
```

MASTER

- Specifies that only the master thread should execute a section of the program

```
#pragma omp master  
structured-block
```

```
int main( )
{
    int a[5], i;
    #pragma omp parallel
    {
        #pragma omp for
        for (i = 0; i < 5; i++)
            a[i] = i * i;

        #pragma omp master
        for (i = 0; i < 5; i++)
            printf("a[%d] = %d\n", i, a[i]);

        #pragma omp barrier

        #pragma omp for
        for (i = 0; i < 5; i++)
            a[i] += i;
    }
}
```


Nested parallelism

```
#pragma omp parallel for
for(j=0; j<jmax; j++){
    #pragma omp parallel for
    for(i=0; i<imax; i++){
        do_work(i,j);
    }
}
```

- OpenMP standard allows but does not require nested parallelism
- Logical function `omp_get_nested()` returns “true” or “false” (1 or 0) to indicate whether or not nested parallelism is enabled in the current region

```
#pragma omp parallel for
for(int y=0; y<25; ++y) {
    #pragma omp parallel for
    for(int x=0; x<80; ++x) {
        tick(x,y);
    }
}
```

- How many ticks will be triggered?

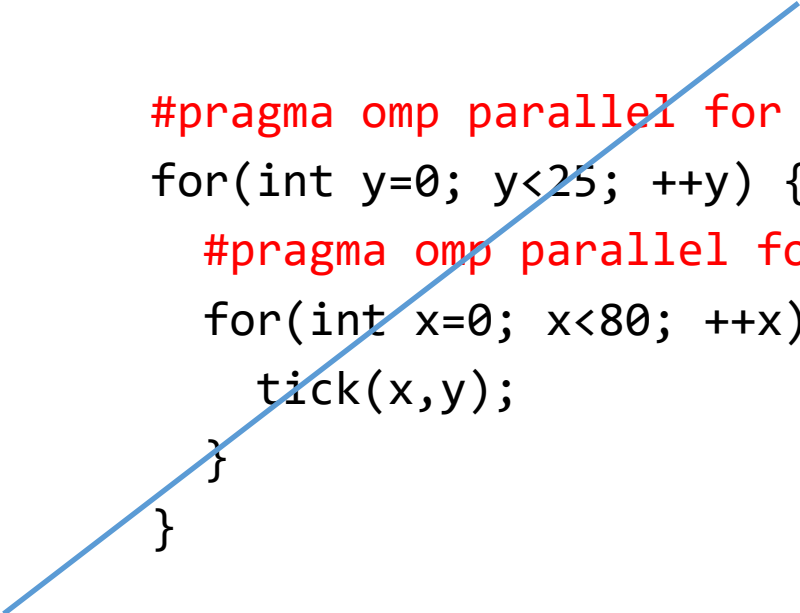
```
#pragma omp parallel for
for(int y=0; y<25; ++y) {
    #pragma omp parallel for
    for(int x=0; x<80; ++x) {
        tick(x,y);
    }
}
```

- How many ticks will be triggered?
- The inner loop is not actually parallelized. Only the outer loop is.
- The inner loop runs in a pure sequence, as if the whole inner `#pragma` was omitted.
- The OpenMP detects that there already exists a team, and instead of a new team of N threads, it will create a team consisting of only the calling thread

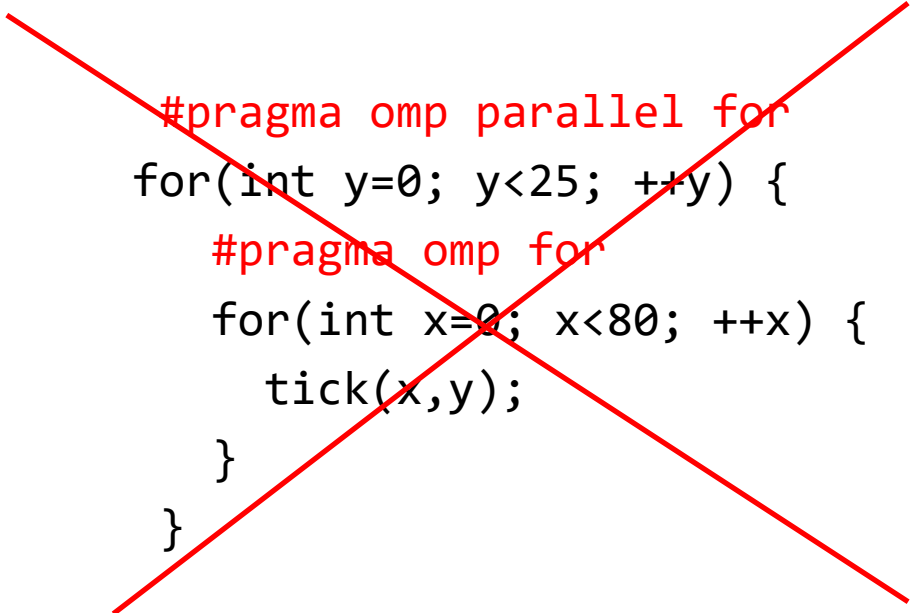
```
#pragma omp parallel for
for(int y=0; y<25; ++y) {
    #pragma omp parallel for
    for(int x=0; x<80; ++x) {
        tick(x,y);
    }
}
```

```
#pragma omp parallel for
for(int y=0; y<25; ++y) {
    #pragma omp for
    for(int x=0; x<80; ++x) {
        tick(x,y);
    }
}
```

- How many ticks will be triggered?
- The inner loop is not actually parallelized. Only the outer loop is.
- The inner loop runs in a pure sequence, as if the whole inner `#pragma` was omitted.
- The OpenMP detects that there already exists a team, and instead of a new team of N threads, it will create a team consisting of only the calling thread



```
#pragma omp parallel for
for(int y=0; y<25; ++y) {
    #pragma omp parallel for
    for(int x=0; x<80; ++x) {
        tick(x,y);
    }
}
```



```
#pragma omp parallel for
for(int y=0; y<25; ++y) {
    #pragma omp for
    for(int x=0; x<80; ++x) {
        tick(x,y);
    }
}
```

- How many ticks will be triggered?
- The inner loop is not actually parallelized. Only the outer loop is.
- The inner loop runs in a pure sequence, as if the whole inner `#pragma` was omitted.
- The OpenMP detects that there already exists a team, and instead of a new team of N threads, it will create a team consisting of only the calling thread

```
#pragma omp parallel for collapse(2)
for(int y=0; y<25; ++y)
    for(int x=0; x<80; ++x)
    {
        tick(x,y);
    }
```

- The number specified in the *collapse* clause is the number of nested loops that are subject to the work-sharing semantics of the OpenMP for construct.

```
#include <stdio.h>
#include "mpi.h"
#include <omp.h>
```

What will this program print if
mpirun -n 2 ./hybrid

```
int main(int argc, char *argv[]) {
    int numprocs, rank, namelen;
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int iam = 0, np = 1;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    #pragma omp parallel default(shared) private(iam, np)
    {
        np = omp_get_num_threads();
        iam = omp_get_thread_num();
        printf("Hello from thread %d out of %d from process %d out of %d\n", iam, np, rank,
numprocs);
    }
    MPI_Finalize(); }
```

```
#include <stdio.h>
#include "mpi.h"
#include <omp.h>

int main(int argc, char *argv[]) {
    int numprocs, rank, namelen;
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    int iam = 0, np = 1;
```

```
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
```

```
    #pragma omp parallel default(shared) private(iam, np)
    {
        np = omp_get_num_threads();
        iam = omp_get_thread_num();
        printf("Hello from thread %d out of %d from process %d out of %d\n", iam, np, rank,
numprocs);
    }
    MPI_Finalize(); }
```

Hello from thread 0 out of 4 from process 0 out of 2
Hello from thread 1 out of 4 from process 0 out of 2
Hello from thread 2 out of 4 from process 0 out of 2
Hello from thread 3 out of 4 from process 0 out of 2
Hello from thread 0 out of 4 from process 1 out of 2
Hello from thread 3 out of 4 from process 1 out of 2
Hello from thread 1 out of 4 from process 1 out of 2
Hello from thread 2 out of 4 from process 1 out of 2


```
#include <omp.h>
#include "mpi.h"
#include <stdio.h>
#define _NUM_THREADS 4
```

What will this program print if
mpirun -n 5 ./hybrid

```
int main (int argc, char *argv[]) {
    int p,my_rank,c;
    omp_set_num_threads(_NUM_THREADS); /* set number of threads to spawn */
    MPI_Init(&argc, &argv); /* initialize MPI stuff */
    MPI_Comm_size(MPI_COMM_WORLD,&p);
    MPI_Comm_rank(MPI_COMM_WORLD,&my_rank);
    #pragma omp parallel reduction(+:c) {
        #pragma omp master {
            if ( 0 == my_rank)
                c = 1;
            else
                c = 2;    } }
    printf("%d\n",c);
    MPI_Finalize();
    return 0;
}
```

```
#include <omp.h>
#include "mpi.h"
#include <stdio.h>
#define _NUM_THREADS 4

int main (int argc, char *argv[]) {
    int p,my_rank,c;
    omp_set_num_threads(_NUM_THREADS); /* set number of threads to spawn */
    MPI_Init(&argc, &argv); /* initialize MPI stuff */
    MPI_Comm_size(MPI_COMM_WORLD,&p);
    MPI_Comm_rank(MPI_COMM_WORLD,&my_rank);
    #pragma omp parallel reduction(+:c) {
        #pragma omp master {
            if ( 0 == my_rank)
                c = 1;
            else
                c = 2;    } }
    printf("%d\n",c);
    MPI_Finalize();
    return 0;
}
```

Concluding Remarks (1)

- OpenMP is a standard for programming shared-memory systems.
- OpenMP uses both special functions and preprocessor directives called pragmas.
- OpenMP programs start multiple threads rather than multiple processes.
- Many OpenMP directives can be modified by clauses.

Concluding Remarks (2)

- A major problem in the development of shared memory programs is the possibility of race conditions.
- OpenMP provides several mechanisms for insuring mutual exclusion in critical sections.
 - Critical directives
 - Named critical directives
 - Atomic directives
 - Simple locks

Concluding Remarks (3)

- By default most systems use a block-partitioning of the iterations in a parallelized for loop.
- OpenMP offers a variety of scheduling options.
- In OpenMP the scope of a variable is the collection of threads to which the variable is accessible.

Concluding Remarks (4)

- A reduction is a computation that repeatedly applies the same reduction operator to a sequence of operands in order to get a single result.